

dt v25.02 Release Notes

- [1. Overview](#)
- 4. New I/O Behaviors
 - [4.1. dtapp I/O Behavior](#)
 - [4.2. hammer I/O Behavior](#)
 - [4.2.1. Hammer Help](#)
 - [4.2.2. Hammer Workloads](#)
 - [4.2.3. Example Hammer Workload](#)
 - [4.3. sio I/O Behavior](#)
 - [4.3.1. sio Help](#)
 - [4.3.2. Example sio Workload](#)

1. Overview

This document describes the release notes for changes introduced in dt version 25.02:

This is the released versions of dt:

```
root@rtp-smc-qa18-4 ~# dt version
```

```
--> Date: December 3rd, 2021, Version: 25.02, Author: Robin T. Miller <--
```

```
root@rtp-smc-qa18-4 ~#
```

[Note: Sections 2 and 3 removed since they are vendor unique!]

4. New I/O Behaviors

My former colleague at NetApp worked with their legal eagles to allow other dt I/O behaviors to be available as open source!

The additional I/O behaviors are: **dtapp**, **hammer**, and **sio**

This means there are five (5) different tools integrated into dt now:

I/O Behaviors:	
iobehavior=type	Specify the I/O behavior. (alias: iob=)
where type is:	
dt	The dt I/O behavior (default).
dtapp	The dtapp I/O behavior.
hammer	The hammer I/O behavior.
sio	The simple I/O (sio) behavior.
thumper	The thumper I/O behavior.

4.1. dtapp I/O Behavior

The dtapp I/O behavior was designed to emulate application I/O across a set of disks. Furthermore, those disks can be mirrored (synchronous or asynchronous). In the case of synchronous mirroring, the data is verified immediately, while in the case of asynchronous mirroring, the verification is done later after async data has been written.

Here's a dt regression test script which shows this behavior. I'm using different paths to the same disk for this example. A list of disks are provided, then dtapp will randomly choose a disk to write to, then verify this data against the mirror disk. Note: This dtapp.dt script is in the dt Scripts/ directory.

```
root@rtp-smc-qa18-4 /usr/local/bin/dt.v25/Scripts# export WRITER="/dev/sdl,/dev/sdm,/dev/sdn"
root@rtp-smc-qa18-4 /usr/local/bin/dt.v25/Scripts# export
MIRROR="/dev/sdam,/dev/sdan,/dev/sdao"
root@rtp-smc-qa18-4 /usr/local/bin/dt.v25/Scripts# dt script=dtapp
dt> version
--> Date: November 23rd, 2021, Version: 25.01, Author: Robin T. Miller <--
dt>
dt> logprefix="%et %prog %device (j:%job t:%thread): "
dt>
dt> $CAPACITYP=capacity=10
dt> $LOGDIR=/var/tmp
dt> system rm -f /var/tmp/dtapp_job*.log /var/tmp/dtapp_thread*.log
dt>
dt> $MIRROR_OPTIONS="iobehavior=dtapp bs=random capacity=10 onerr=abort history=25
keepalivet=30 tag=mirror job_log=/var/tmp/dtapp_job%job.log log=/var/tmp/dtapp_thread-
j%jobt%thread.log"
dt>
dt> $WRITER=/dev/sdl,/dev/sdm,/dev/sdn
dt> $MIRROR=/dev/sdam,/dev/sdan,/dev/sdao
dt>
dt> of=/dev/sdl,/dev/sdm,/dev/sdn mirror=/dev/sdam,/dev/sdan,/dev/sdao iobehavior=dtapp bs=random
capacity=10 onerr=abort history=25 keepalivet=30 tag=mirror job_log=/var/tmp/dtapp_job%job.log
log=/var/tmp/dtapp_thread-j%jobt%thread.log limit=250m
```

4.2. hammer I/O Behavior

The *hammer* I/O behavior is a file system tool. Hammer performs file system operations, including file locking, delete, rename, and truncate, in addition to writing and reading files. An in-memory list of files is maintained and randomly chosen. By default, hammer runs forever and also supports file system full and file op percentages.

4.2.1. Hammer Help

```
root@rtp-smc-qa18-4 /# dt iobehavior=hammer help
```

```
Usage: dt iobehavior=hammer [dir=DirectoryPath] [options...]
```

Options:

-help	Show this help text, then exit.
dir=DirectoryPath	Directory path for hammer files.
	If omitted, current directory is the default.

file). -api posix -bg -checkinodes -filercore=FILENAME -nofilercore -direct -fill -interactive -iterations NUMBER -logfile FILE -mode={mixed creates overwrites} (Default: mixed) -nocleanup -noflush -onlydelete -onlyflush -runtime=N -threads=value -seed=value -version	The only API supported for Unix is POSIX. Don't use stdin or stdout (output to log File system reusing inodes? Coredump the filer upon corruption. Disable coredumping the filer. Disable filesystem caching. Fill disk and then keep it full. Use interactive mode (not supported). The number of iterations to execute. Use logfile FILE. (Default is none). Use the specified mode. Don't remove files upon completion. Don't use flush buffers (write through). When exiting, only delete files. Always use async I/O and flush buffers. Stop hammering after N seconds. The number of hammer threads. Set the random seed to use. Print the version, then exit.
--	--

Blocksize Options:

-blocksize=NBYES -blocksize=random -minbsize=NBYES 512) -maxbsize=NBYES 131072) -minfsize=NBYES 1) -maxfsize=NBYES 5242880)	Use blocksize NBYES. Use a random blocksize. (Default: True) Set minimum block size to NBYES. (Default: Set maximum block size to NBYES. (Default: Set minimum file size to NBYES. (Default: Set maximum file size to NBYES. (Default:
--	---

Error Control Options:

-ignorelockerrors -ignorefileerrors continue. -ignoredatacorruption continue. -ignoreallerrors continue. errors. properly].	Don't halt on file locking errors, continue. Don't halt on file operation errors, Don't halt on data corruption errors, Don't halt on any of the above errors, [NOTE: hammer will stop on other critical that can prevent it from functioning
---	--

Lock Control Options:

-nolockdebug chatty). -lockfiles defaults for the lock options below. -lockmode={mixed full partial} (default: mixed). -unlockchance=[0-100] unlocking, 0-100 percent.	Exclude file lock/unlock debug output (it's Include file locks (locks & unlocks) using More chance of full or partial file locks Probability of keeping locks and skipping
---	---

```

Examples:
    if -unlockchance=100    100% chance of unlocking, ALL files
unlocked. [default]
    if -unlockchance=50     50% chance of unlocking each file.
    if -unlockchance=0      0% chance of unlocking, NO files are
unlocked.

```

dt Options Supported:

```

    bufmodes={buffered,unbuffered,cachereads,cachewrites}
                                Set one or more buffering modes (Default:
none)
    maxdata=value                The maximum data limit (all files).
    maxdatap=value              The maximum data percentage (range: 0-100).
    maxfiles=value              The maximum files for all directories.
    stopon=filename              Watch for file existence, then stop.
    enable=raw                   The read after write flag.

```

Also know, I/O monitoring (noprogram options), keepalive, and trigger= options are also supported with hammer.

Examples:

```

% dt iobehavior=hammer dir=/mnt/hammer maxdatap=25 runtime=1h
% dt iobehavior=hammer dir=/mnt/hammer bufmodes=buffered,unbuffered
stopon=stopfile
% dt iobehavior=hammer dir=/mnt/hammer -lockfiles -onlydelete -threads=3
log=hammer.log

root@rtp-smc-qa18-4 /#

```

4.2.2. Hammer Workloads

These are the predefined *hammer* workloads:

```

hammer: Standard hammer workload
iobehavior=hammer

hammer_bufmodes: Hammer workload using buffered/unbuffered modes
iobehavior=hammer bufmodes=buffered,unbuffered

hammer_locking: Hammer workload with file logging and stop file
iobehavior=hammer -lockfiles -onlydelete stopon=/var/tmp/stopdt

```

4.2.3. Example Hammer Workload

```

dt (j:1 t:1): Command Line:
dt (j:1 t:1):
dt (j:1 t:1): # ./dt of=/mnt/localhost/iscsi-rtp-smc-qa18-4-v7-4421cc0f0efae5826c9ce900fa42bd6b
workload=hammer sdsf=/dev/mapper/mpathbo logdir=/var/tmp/dtlogs_20211123-142143-300785 job_log=dt-
iscsi-rtp-smc-qa18-4-v7-4421cc0f0efae5826c9ce900fa42bd6b-j%job.log log=dt-iscsi-rtp-smc-qa18-4-v7-
4421cc0f0efae5826c9ce900fa42bd6b-j%jobt%thread.log enable=async onerr=abort stopon=/var/tmp/stopdt
tag=iotag runtime=1m
dt (j:1 t:1):
dt (j:1 t:1): --> Date: November 18th, 2021, Version: 25.01, Author: Robin T. Miller <--

```

```

dt (j:1 t:1):
dt (j:1 t:1): Copyright (c) 2012 Network Appliance, Inc. All rights reserved.
dt (j:1 t:1): hammer started at Tue Nov 23 14:21:45 2021
dt (j:1 t:1): version=$Id: hammer.c#10 $
dt (j:1 t:1): path=/mnt/localhost/iscsi-rtp-smc-qa18-4-v7-
4421cc0f0efae5826c9ce900fa42bd6b/52b2c4d5367e759fd296b049c8cc7e9a pid=0x6D49
dt (j:1 t:1): client=rtp-smc-qa18-4 (Linux 3.10.0-1160.2.1.el7.x86_64 #1 SMP Mon Sep 21 21:00:09 EDT 2020
x86_64)
dt (j:1 t:1): minfsize=0x1 maxfsize=0x500000
dt (j:1 t:1): minbsize=0x00000200 maxbsize=0x00020000 blocksize=random
dt (j:1 t:1): api=posix mode=mixed streams=off flush=random nocleanup=false retrydisc=true
seed=7033918714584814846
dt (j:1 t:1): logfile=/var/tmp/dtlogs_20211123-142143-300785/dt-iscsi-rtp-smc-qa18-4-v7-
4421cc0f0efae5826c9ce900fa42bd6b-j1t1.log timezone=EST
dt (j:1 t:1): num_iterations=0 max_iterations=18446744073709551615 cur_runtime=0 max_runtime=60
dt (j:1 t:1): 2021/11/23-14:21:45 CREATE 00000001.ham fileid=0x6B8B4567 blocksize=0x00019E7D
filesize=0x0043503b timestamp=0x619D3F49 (sync) 70992.8K/s
dt (j:1 t:1): 2021/11/23-14:21:45 TRUNC 00000001.ham fileid=0x6B8B4567 oldsize=0x0043503b
newsize=0x000f0f42 0.000987sec
dt (j:1 t:1): 2021/11/23-14:21:45 CREATE 00000002.ham fileid=0x327B23C6 blocksize=0x000061CD
filesize=0x00244c28 timestamp=0x619D3F49 (sync) 28500.7K/s
dt (j:1 t:1): 2021/11/23-14:21:46 CREATE 00000003.ham fileid=0x643C9869 blocksize=0x00002B87
filesize=0x00062eca timestamp=0x619D3F4A (sync) 13948.2K/s
dt (j:1 t:1): 2021/11/23-14:21:46 TRUNC 00000002.ham fileid=0x327B23C6 oldsize=0x00244c28
newsize=0x00101ca7 0.000356sec
dt (j:1 t:1): 2021/11/23-14:21:46 CREATE 00000004.ham fileid=0x66334873 blocksize=0x00010927
filesize=0x0015d12a timestamp=0x619D3F4A (async then flush) 262411K/s
dt (j:1 t:1): 2021/11/23-14:21:46 TRUNC 00000001.ham fileid=0x6B8B4567 oldsize=0x000f0f42
newsize=0x00083dda 0.000167sec
dt (j:1 t:1): 2021/11/23-14:21:46 TRUNC 00000001.ham fileid=0x6B8B4567 oldsize=0x00083dda
newsize=0x00078f0e 4e-05sec
dt (j:1 t:1): 2021/11/23-14:21:46 READ 00000001.ham fileid=0x6B8B4567 blocksize=0x00014314
filesize=0x00078f0e 842794K/s
dt (j:1 t:1): 2021/11/23-14:21:46 DELETE 00000001.ham fileid=0x6B8B4567 0.000178sec
dt (j:1 t:1): 2021/11/23-14:21:46 TRUNC 00000003.ham fileid=0x643C9869 oldsize=0x00062eca
newsize=0x00040804 6.3e-05sec
dt (j:1 t:1): 2021/11/23-14:21:46 RENAME 00000003.ham fileid=0x643C9869 newpath=00000005.ham 3.3e-05sec

```

4.3. sio I/O Behavior

The *sio* I/O behavior is a performance tool. As a performance tool, no data verification is performed by default. Performance statistics are reported upon completion.

4.3.1. sio Help

```
root@rtp-smc-qa18-4 /# dt iobehavior=sio help
```

```
Usage: dt iobehavior=sio [options...]
```

sio (Simple I/O Load Generator) - NetApp

A tool to generate artificial I/O workloads against any device
 Supports numerous configuration variables (reads vs writes, etc)
 Supports multiple devices and multiple threads. Collects a wide
 variety of statistics on I/O client machines and/or I/O servers.

Basic Usage:

```
dt iobehavior=sio readp=<read%> randp=<rand%> bs=<blksize> starting=<start> \
    end=<end> runtime=<secs> threads=<threads> devs=<dev>,[devs,...]
readp=<read>          Percentage of accesses that are reads. Range [0,100].
                      'random' keyword makes the read/write percentage
random.              BEWARE, writing to a file is unchecked and will trash
                      files.
randp=<rand>          Percentage of accesses that are random. Range [0,100].
                      Sequential accesses = 0%, else random percentage
bs=<blksize>          Size of I/O's. Example: 2k, 4k, 1m
                      'random' keyword makes the I/O size random 512 bytes to
262144 bytes.
ibs=<blksize>          Size of read requests. (overrides bs= option)
obs=<blksize>          Size of write requests.
start=<start_byte>     Lower bound for access location in each file.
end=<file_size>        Total bytes accessed in each file (e.g. 100m, 2g,
1000k).
runtime=<seconds>      Runtime for test. Counting starts AFTER all threads
have started.
threads=<numthreads>   Concurrent I/O generators. Uses real individual
threads.
dev=<dev>              Device to access. May be file (foo.out) or device
(/dev/dsk/etc).
    or devs=<dev>[,...] Multiple devices and/or files can be specified, comma
separated.
    or file=<paths>      One or more paths to files to access (synonym for
'devs' option).
```

Examples:

```
1) Random 4k I/O with 25% reads/75% writes, 75% random/sequential for 10
minutes.
    Accessing a total of 250 megabytes in each file, after prefilling the
    file.
    % dt iobehavior=sio file=a.file,b.file bs=4k readp=25 randp=75 end=250m -
prefill runtime=10m
2) Random reads and writes with random block sizes via 10 threads to the
same file.
    This test will run infinitely without -numops or runtime options.
    % dt iobehavior=sio file=a.file bs=random readp=random end=1g -direct -
verify threads=100
```

Options:

```
Options are divided into four categories:
    Basic Features, Advanced Features,
    Q/A Features, and Esoteric Stuff.
```

Basic Features:

```
-help          List this sio help, then exit.
-version       Display this detailed version log.
-noflock       Do NOT lock files. Locking affects caching on some
OS's.
-noheader      Suppress single line header output. (Good for
multiple runs).
-debug         Output detailed debug info. Be prepared for a lot
of info.
```

-Debug	Very verbose debug information. Be prepared for a lot of info!
-niceoutput	Print output in single column, human-readable format.
-prettyprint	Pretty print the output (this is dt's format).
-no_dsync	Do NOT open files with O_DSYNC. Allows async writes.
-noperf	Do NOT display performance statistics.

Advanced Features:

-stop=<fname>	Watch for existence of file 'fname' and terminate.
-think=<msec>	Each thread waits 'ms' MS before issuing each I/O.
-iops=<rate>	Target IOPS for each thread.
-lockall	Lock the complete file as opposed to a single byte.
-truncate:	IFF pure sequential writes, then when I/O wraps to beginning of file, the file is truncated.
-max_blksize=<bytes>	Set maximum block size to 'bytes'.
-max_latency=(ms)	Maximum allowed latency (in milliseconds) of an IO.
-direct	Disable filesystem caching.
-align=(size)	Alignment to be used with random block size.
-break_on_dc	Exit upon detecting data corruption ASAP.

Q/A Features:

-verify	Read back written data and verify content.
-verify_retry=<n>	Retry failed verifies 'n' times.
-instrumentation	Special pattern insertion technique.
-fixedfill=<value>	Fill the file with 8 bit value.

Esoteric Stuff:

-numops=<num_ops>	Run for 'num_ops' I/O's and stop. Beware stats.
-fileperthread	Open one file per thread. Special names.
-blockno	Prints out the I/O block numbers.
-iofailok	Allow I/O failures (do not access file again).
-iomutex	Use mutex to synchronize multiple threads.
-fillonce	Write all files once, then stop.
-prefill	Write all files prior to test I/O.
-partition_among_threads	Partition the file among threads.

root@rtp-smc-qa18-4 /#

sio Workloads

These are the predefined *sio* workloads:

```
root@rtp-smc-qa18-4 /# dt workloads sio
Valid workloads:
```

```
    sio_percentages: sio workload with percentages, 4k I/O, 1g file, and prefilling
                      iobehavior=sio bs=4k readp=25 randp=75 end=1g -prefill runtime=1h

    sio_random: sio workload with random block sizes, random read/write, 10g file w/verification
                 iobehavior=sio bs=random readp=random end=10g -direct -verify threads=10

    sio_many_files: Populate directory with many files concurrently (write only)
                    iobehavior=sio bs=32k end=10m files=100 passes=1

    sio_many_slices: Populate device or file via many slices (write only)
                     iobehavior=sio bs=32k end=10g slices=100 passes=1

    sio_many_threads: Populate device or file via many threads (write only)
                      iobehavior=sio bs=32k end=10g threads=100 passes=1
```

```
root@rtp-smc-qa18-4 /#
```

4.3.2. Example sio Workload

```
root@rtp-smc-qa18-4 /var/tmp/dt/logs# dt of=/mnt/localhost/iscsi-rtp-smc-qa18-4-v7-4421cc0f0efae5826c9ce900fa42bd6b/sio.data sdsf=/dev/mapper/mpathbo workload=sio_many_files end=1g
dt (j:1 t:0): Read: 0 Rand: 0 BlkSz: 32768 BegnBlk: 0 EndBlk: 32768 Secs: -1 Threads: 100 Devs: 1 /mnt/localhost/iscsi-rtp-smc-qa18-4-v7-4421cc0f0efae5826c9ce900fa42bd6b/sio.data
dt (j:1 t:1):
dt (j:1 t:1): Thread Latency Stats:
dt (j:1 t:1): Thread: 1
dt (j:1 t:1): ios: 32768
dt (j:1 t:1): latency(us): 186980130
dt (j:1 t:1): sumofsquares: 110234
dt (j:1 t:1): min(ms): 0.41
dt (j:1 t:1): max(ms): 88.36
dt (j:1 t:1): avg(ms): 5.71
dt (j:1 t:1): stddev(ms): 1.83
dt (j:1 t:2):
dt (j:1 t:2): Thread: 2
dt (j:1 t:2): ios: 32768
dt (j:1 t:2): latency(us): 186975277
dt (j:1 t:2): sumofsquares: 112283
dt (j:1 t:2): min(ms): 0.41
dt (j:1 t:2): max(ms): 73.88
dt (j:1 t:2): avg(ms): 5.71
dt (j:1 t:2): stddev(ms): 1.85
```

Full output emitted.

```
dt (j:1 t:100):
dt (j:1 t:100): Thread: 100
dt (j:1 t:100): ios: 320
dt (j:1 t:100): latency(us): 1734615
dt (j:1 t:100): sumofsquares: 145
dt (j:1 t:100): min(ms): 2.90
dt (j:1 t:100): max(ms): 10.19
dt (j:1 t:100): avg(ms): 5.42
dt (j:1 t:100): stddev(ms): 0.67
dt (j:1 t:1):
dt (j:1 t:1): IOPS,TPUT(KB/s),LAT(ms)Calc'd,LAT(ms)Actual,READ,RAND,IOS,SEC,THDS,BLKSZ
dt (j:1 t:1): 16000,512000,6.250,[5.410],0,0,32000,2,100,32768
dt (j:1 t:1): Global Latency Stats:
dt (j:1 t:1): ios: 32000
dt (j:1 t:1): latency(us): 173132086
dt (j:1 t:1): sumofsquares: 11071
dt (j:1 t:1): min(ms): 1.17
dt (j:1 t:1): max(ms): 10.66
dt (j:1 t:1): avg(ms): 5.41
dt (j:1 t:1): stddev: 0.59
dt (j:1 t:1):
dt (j:1 t:1): global_reads = 0; global_bytes_read = 0 (0 KB)
dt (j:1 t:1): global_writes = 32000; global_bytes_written = 1048576000 (1024000 KB)
dt (j:1 t:1): global_time_start = 1637698093, global_stop_time = 1637698095
dt (j:1 t:1): measurement start = 1637698093, measurement stop = 1637698095
dt (j:1 t:1): Computed run time seconds = 2
dt (j:1 t:1): Computed measurement seconds = 2
dt (j:1 t:1): Computed IOPS = 16000.00
dt (j:1 t:1): Computed KB/s = 512000.00
dt (j:1 t:1): Computed bytes/IO = 32768
```